

Active Learning for Drug Selection on Identified Target Protein

Christine Baek
christib@andrew.cmu.edu

Qi Chu
qchu@andrew.cmu.edu

1 INTRODUCTION

In this project, we explore three datasets of different noise and/or labels, for identification of compounds that bind to specific target protein. We first perform feature selection using χ^2 test, then adopt the *largest positive score* [1] for active learning strategy, and SVM for model learning. We also perform a gradient experiment to learn the optimal balance for allocating queries between feature selection and model learning.

2 METHODS

Input :

- Training data consists of 4000 training instances, each with 1000 features and true label
- Test data consists of 1000 testing instances, each with 1000 features and true label
- Blind prediction data consists of 1000 instances, each with 1000 features
- Each data(row) represents potential compound that may bind to the target
- Each feature(column) represents structural or chemical

As with real world drug target screening, we began with assumption that only an extremely small fraction of compounds will bind to the target. Also, only a small fraction of the features would be relevant in predicting the label (are actual sites that the drug can potentially bind to).

Our challenge is to first find which features are relevant in accurately predicting whether the compound will bind to the target or not, and then learning on those features.

We performed gradient experiment, and measured the performance based on the number of queries made between initial learning and active learning, as the total number of queries is capped at 2500. Our predictions are then compared against random learner, as well as 0-guess, which predicts inactivity (label=0) for all compounds.

2.1 LEARNING STRATEGY

Because of the imbalanced dataset (only small portion of the molecules actually bind to target), it would have been dangerous to use active learning strategy without any prior. We decided to split our learning into two different phases : Initial and Active. During initial learning, we randomly select molecules without any assumptions or selection criteria, other than choosing molecules we have not seen before. We perform feature selection during this phase. Once initial learning is over, we begin the active learning phase where points are selected per *largest positive* strategy as discussed in [1], elaborated below. Total number of queries made during both initial and active phase total 2500, the allotted budget.

2.1.1 INITIAL FEATURE SELECTION

Randomly chosen molecules are queried for their label and features. This initial phase is important in learning which features are relevant, in order to build a more relevant model. While data has been omitted, our previous attempts without explicit feature selection has returned low F1 score as well as error rate.

2.1.2 ACTIVE COMPOUND SELECTION

We used selection strategy of *largest positive score* which performed the best in [1]. Given that most compounds are inactive and would be predicted to be 0, this strategy searches for the compounds that are active and effective in discovering potentially active compounds. Given our goal of drug discovery for our target, false positives (drugs that are classified as binding to target, but does not) can be tolerated within reason, but false negatives are potentially costly. We do not expend our limited number of queries for looking up inactive compounds, but instead dedicate our efforts towards finding potentially active compounds which are far and few between. *Largest positive* score compound is selected by computing a *score* of activity based on signed distance from the SVM hyperplane. Note that this strategy works for our purpose of finding as many positive molecules that bind to drug target, but a different strategy such as choosing compounds near boundary may be more suited for studying the structure-activity relationship [1]

2.2 CLASSIFIER STRATEGY

TODO: elaborate

SVM was chosen as our classifier for multiple reasons. SVM finds a hyperplane that maximizes the distance between the nearest data (support vectors) and the hyperplane. Given the

high dimension of our input data (1000 features) with assumption of linear separability, it was important that a classifier that scales easily to high dimension, with no local optima. SVM separates inactive from active compounds by using maximum margin hyperplane for decision boundary. Gaussian kernel (default) was used for training.

3 RESULTS

Gradient experiment was performed for varying number of queries divided between the initial learning (feature selection) and active learning (model learning). Graphs for each section (difficulty) include the best performing division of queries, but full data is available

3.1 EASY




Figure 3.1: Error rate of easy test set plotted against number of queries made to oracle




Figure 3.2: F1 Score of easy test set plotted against number of queries made to oracle

3.2 MODERATE




Figure 3.3: Error rate of moderate test set plotted against number of queries made to oracle




Figure 3.4: F1 score of moderate test set plotted against number of queries made to oracle

3.3 DIFFICULT




Figure 3.5: Error rate of difficult test set plotted against number of queries made to oracle




Figure 3.6: F1 score of difficult test set plotted against number of queries made to oracle

Table 3.1: F1 score of various parameters

# of queries - Initial Learning	500	1000	1500	2000
# of queries - Active Learning	2000	1500	1000	500
EASY	0.6409	0.6328	0.9375	0.9375
MODERATE	0.2625	0.4365	0.4686	0.5034
DIFFICULT(pos. v.s. neg.)	0.4286	0.5133	0.5346	0.6481

F1 score for 3 different input data, with gradient of initial vs active learning query division. This table shows the F1 score of each test data after the model had been trained on varying number of queries for Initial Learning (feature selection) and active learning (model learning).

Table 3.2: Error rate of various parameters

# of queries - Initial Learning	500	1000	1500	2000
# of queries - Active Learning	2000	1500	1000	500
EASY	0.065	0.065	0.014	0.014
MODERATE	0.191	0.142	0.161	0.146
DIFFICULT(pos. v.s. neg.)	0.622	0.346	0.375	0.301

Error rate for 3 different input data, with gradient of initial vs active learning query division. This table shows the error rate of each test data after the model had been trained on varying number of queries for Initial Learning (feature selection) and active learning (model learning).

4 CONCLUSION

TODO: briefly summarize

REFERENCES

- [1] Warmuth, Liao, Ratsch, Mathieson, Putta and Lemmen, *Active Learning with Support Vector Machines in the Drug Discovery Process*, J. Chem. Inf. Comput. Sci. 2003, 43, 667-673